MUSEUM OF ARTIFICIAL INTELLIGENCE

A NEW EXHIBITION · 1950–2026

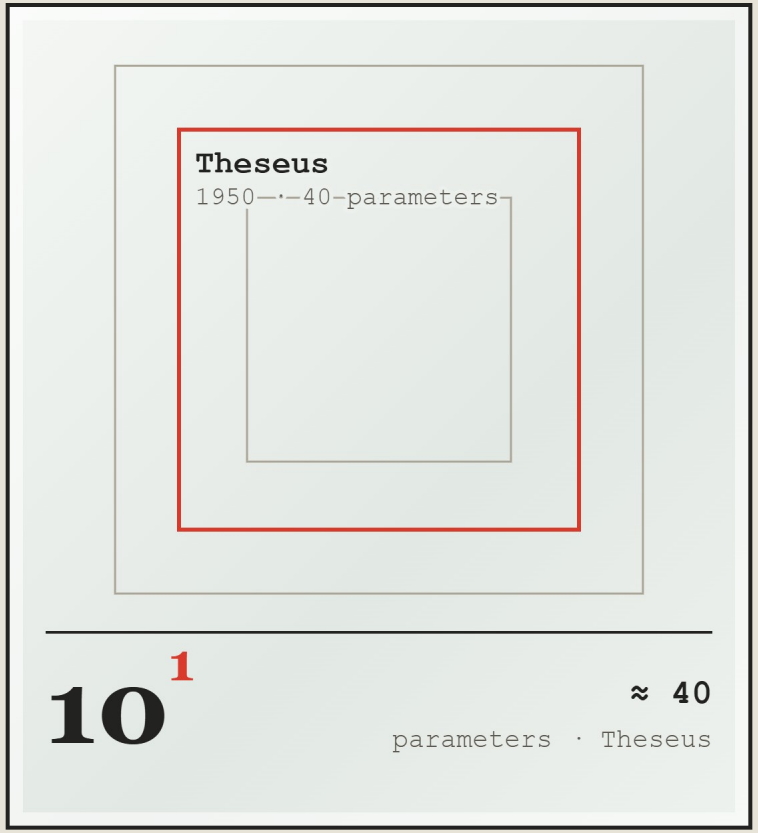
Powers of Ten

How AI models grew from 40 parameters to 100 trillion.

CURATED BY

Lujain Al Ibrahim

Final Project · Unit 3



OBJECT 001 · LIVE EXHIBIT

Zooming out through AI history

Each square is ten times the last.

3,628 AI models

This collection, lent by the research group Epoch AI, records every documented AI model from 1950 to October 2026: who built it, when, how large it is, and how much computing it took to train.

COLLECTION FACTS

Objects	**3,628**
Catalogue fields	**57**
Records left blank	**61%**
Most objects from	**OpenAI, Alibaba, Google**
Why chosen	**AI is in the news daily**

AI.1950.0001 · earliest

Theseus

Claude Shannon's maze-solving robot mouse

40

AI.2022.BGL · largest

BaGuaLu

Trained on a Chinese supercomputer

10^{14}

↑ parameters (model size)

CONSERVATION REPORT

Fields 57 → **22**
Blank cells 61% → **17%**
Domain labels 191 → **22**
Undated records removed: 12

Four questions this exhibition answers

Q.01

How fast have AI models grown?

field: parameters

Q.02

Who builds them: industry or universities?

field: org_type

Q.03

Which domains dominate, once cleaned?

field: domain

Q.04

Can size and year predict a model's cost?

field: compute_flop

ALONG THE CORRIDOR · MILESTONES IN THE COLLECTION

1950 Theseus

1957 Perceptron

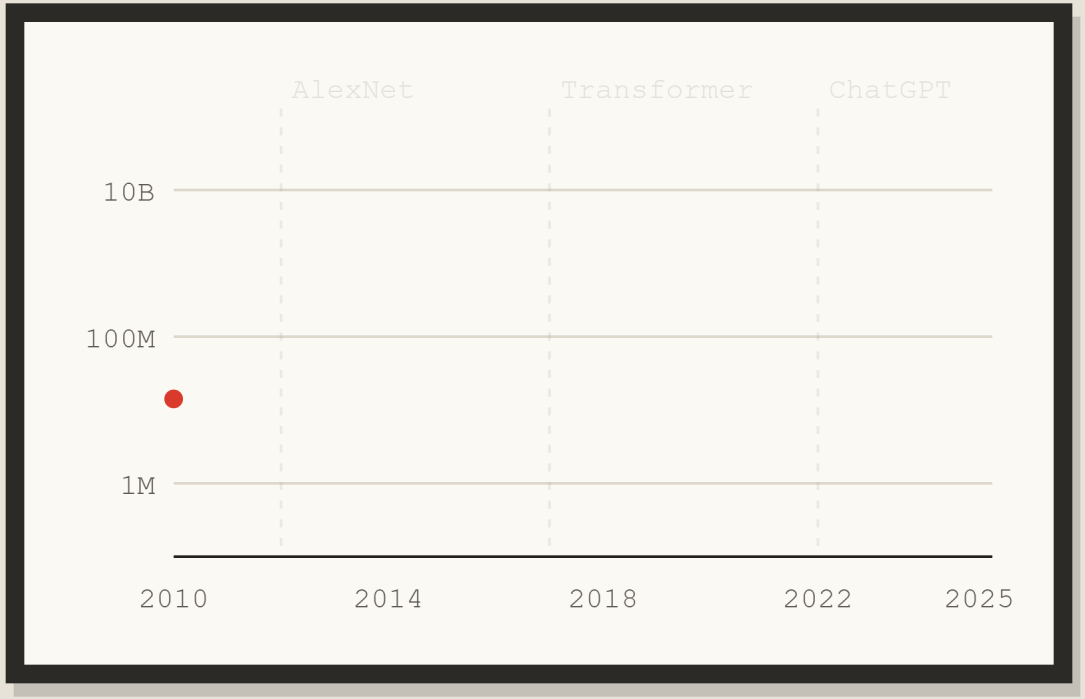
1997 Deep Blue

2012 AlexNet

2017 Transformer

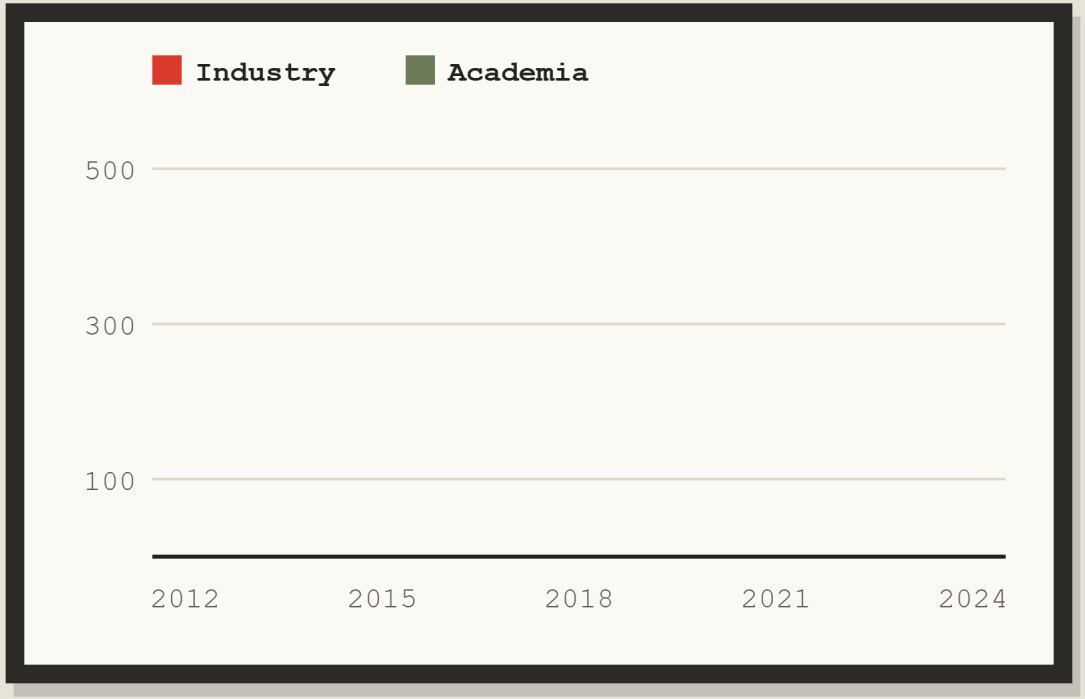
2020 GPT-3

Bigger, faster, and built by industry



Typical model size, 2010–2025

Median parameters per year, log scale.
Live exhibit.

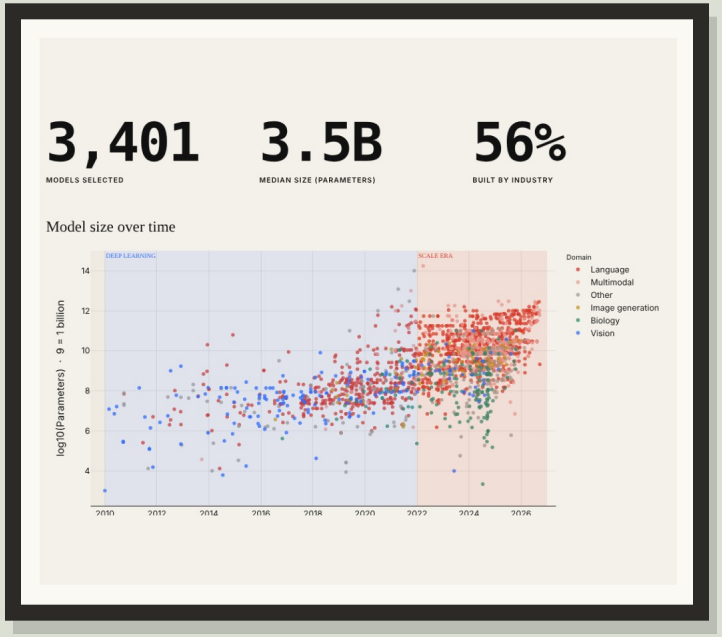


Who builds AI, 2012–2024

Models released per year. Live exhibit.

The dashboard: a story you can explore







The mess, then the clean-up

Real raw rows with their problems circled, then fixed.

Findings you can filter

Year slider, domain tags and a dropdown update every chart.

Predict the compute

Enter a size and year; a regression estimates the cost.

Four discoveries from the archive

EXHIBIT A

12 mo

AI is growing faster than ever

From AlexNet (60 million, 2012) to GPT-3 (175 billion, 2020): size now doubles roughly every 12 months.

EXHIBIT B

70×

Industry builds the biggest models

OpenAI, Alibaba and Google built the most. A typical industry model is 70× a university one.

EXHIBIT C

191→22

Cleaning uncovered a hidden giant

Multimodal models, scattered across 170 spellings, are the second-largest domain.

EXHIBIT D

R^2 .85

A model's cost can be estimated

Size and year alone explain 85% of the differences in training compute, typically within about 11×.

What I learned, and what's missing

LESSONS FROM THE WORK

Cleaning decisions shape the results: leaving gaps empty was more honest than filling them, and merging labels revealed new patterns.

MISSING PIECES

Companies often hide size and compute, so the newest giants are likely underestimated, and 2025–2026 is still incomplete.

FUTURE ACQUISITIONS

Compare countries, track open versus closed models, and try a more flexible model such as a random forest.

VISITOR'S BOOK

AI didn't just grow. It grew by powers of ten.

Questions welcome